

SCS-RG — Documentation

Multi-Hazard Smart Campus Safety & Response Grid. RoboFusion 1.0 Techathon, Round 1. Team: [TEAM_NAME] .
Track B (Wokwi ESP32) primary — see Track declaration.


This file is the source for the submitted PDF (export via any Markdown→PDF tool, e.g. Pandoc or VS Code's "Markdown PDF" extension, after inserting the Wokwi screenshots noted in Circuit Diagrams).

1. System Overview

Three campus zones (IoT Lab, Server Room, Data Science Lab) each run an identical ESP32 node that reads raw hazard sensors and posts them to a central FastAPI backend every 750 ms. The backend — never the node — fuses those readings into a risk score, classifies SAFE/WARNING/CRITICAL, persists every state transition and incident, and is the only thing that ever tells a node to sound its buzzer/relay. A single-page dashboard subscribes over WebSocket and renders exactly what the backend/DB say, never deriving state on its own.

2. Architecture

FIGURE 1 — SYSTEM ARCHITECTURE



Key properties this diagram is meant to make visible:

- **Single source of truth (rule 2):** every REST response and every WS push for zone/priority state is built by exactly one function, `status_service.build_zone_status_payload()`. The dashboard never derives state client-side.
- **Actuation is one-directional and sensor-only (rule 4):** `Commands` reads only `zone_transitions` (which only `pipeline.py`'s sensor/override path writes). `n1_report.py` and the Bonus-3 prediction path (if built) have no code path into `Commands` or `zone_transitions` at all — not "we choose not to call it," there is no import edge.
- **Restart safety (rule 6):** `zone_manager.py`'s in-memory cache is rebuilt from `zone_transitions` / `readings` in `main.py`'s lifespan startup, before the app accepts any connection.

3. Risk Fusion Formula (locked)

```
risk_score(zone) =  
    40 * fire_signal          (0-1 continuous, debounced entry / decaying exit)  
+ 25 * gas_level_norm        (0.0-1.0, MQ-2 normalized; 0 during 30s warm-up)  
+ 25 * water_level_norm      (0.0-1.0)  
+ 10 * occupancy_factor      (0 or 1, PIR after a 1.5s hold)
```

```
SAFE < 30    |    30 <= WARNING < 65    |    CRITICAL >= 65  
Hysteresis: exit CRITICAL only below 55, with a minimum 3s hold.
```

Implemented verbatim in `backend/fusion.py` (the score) and `backend/state_machine.py` (the band + hysteresis) — weights are named constants in `backend/config.py` (`WEIGHT_FIRE=40`, `WEIGHT_GAS=25`, `WEIGHT_WATER=25`, `WEIGHT_OCCUPANCY=10`), not re-derived or duplicated anywhere else, so the numbers in this document and the numbers the backend actually computes with cannot drift apart (TC30).

Justification. Fire is weighted highest because it is the fastest-escalating and most destructive hazard in electronics-dense labs — a debounced flame reading alone (40) sits just below WARNING, but combined with almost any secondary hazard reading it crosses into CRITICAL within one reading cycle. Water is raised to parity with gas (both 25) because two of our three zones are server/GPU rooms, where a coolant or condensate leak is the realistic catastrophic failure mode, not a hypothetical one — a flood-equivalent hazard deserves the same weight as a gas leak, not a token afterthought — and, unlike fire (40) alone, a debounced flame reading by itself already crosses into WARNING; either gas or water raised to parity means a second hazard reliably tips a WARNING zone into CRITICAL within one reading cycle, rather than requiring a third signal. Occupancy is weighted lowest (10) *in a zone's own score*, deliberately: an empty zone on fire is still a full equipment/infrastructure emergency regardless of who's in the room, so occupancy shouldn't be able to suppress a hazard score. Its real weight belongs in **cross-zone triage** — deciding which of several simultaneously-CRITICAL zones responders reach first — where life safety should dominate. That's why occupancy reappears, much more heavily, in the priority formula below (life > assets) instead of being duplicated in the per-zone score.

Sensor pipeline parameters (all named constants in `backend/config.py`, enforced in `backend/zone_manager.py`):

Parameter	Value	Effect
Sampling interval	750 ms	Node POST cadence
Flame debounce	5 consecutive HIGH readings	Brief flicker never triggers
Flame decay	linear to 0 over 5 s	No instant snap on removal
Gas warm-up	ignored for 30 s after boot (<code>uptime_ms</code>)	No false trigger on cold start
PIR hold	1.5 s before a state change is applied/logged	No log/score spam on flicker
Offline threshold	no reading for > 3 s (4× interval)	Distinct OFFLINE badge, never SAFE

4. State Machine

`backend/state_machine.py::classify()` — pure function, no I/O, unit-tested directly. CRITICAL is sticky: once entered, the zone cannot exit until both conditions hold simultaneously — score has dropped below 55 (not 65, to avoid immediate flip-flop right at the boundary) **and** at least 3 s have elapsed since CRITICAL was entered. SAFE and WARNING have no hysteresis; they reclassify on every reading purely from the current score.

5. Priority Ranking (locked, cross-zone triage)

```
priority = risk_score
    + 15 * occupancy_factor           (people present -> jump the queue)
    + min(10, unacked_seconds / 15)  (escalates while ignored, cap +10)
    + advisory_boost (Bonus 4, cap +10, decays over 10 min)
Tie-break: earlier CRITICAL entry first.
```

Computed only over zones **currently CRITICAL** (`backend/priority.py`). Occupancy dominates here rather than in the per-zone score because this is exactly the "which fire do we run to first" decision, where life safety should outweigh raw hazard magnitude. The `occupancy_factor` used *here* is PIR backed up by the Bonus 1 camera cross-check (`PIR OR fresh camera detection` — see §11); the zone's own risk score stays PIR-only per the locked formula. `unacked_seconds` escalates a zone's priority the longer it sits un-acknowledged, capped at +10 so an old, low-risk incident can never permanently outrank a fresh, severe one. `advisory_boost` (Bonus 4) is the only non-sensor-derived term in the system, and it is deliberately confined to *this* formula — see §10 NL Incident Reporting.

The dashboard renders the exact breakdown per zone, e.g.:


```
Risk 78 + Occupied +15 + unacked 90s +6 = 99
```

so a judge can verify the ranking without trusting a black box (TC12c).

6. Database Schema

SQLite via SQLAlchemy, WAL journal mode, `PRAGMA foreign_keys=ON` (see `backend/database.py`). All foreign keys are `ON DELETE RESTRICT` — a zone with any historical row (a sensor, a reading, a transition, an incident) cannot be deleted, only left in place (TC18b).

FIGURE 2 — ENTITY RELATIONSHIPS



Indexes: `idx_incidents_status_created(status, opened_at)` — the `GET /api/incidents?status=&from=&to=` filter path (TC19); verified to stay fast (~14 ms) against 300+ seeded incidents / 10k+ seeded readings via `sim/seed.py --readings 10000 --incidents 300`. `idx_readings_zone_ts(zone_id, ts_server)` — per-zone chronological reads (trend window, timeline queries).

Three additive columns beyond the literal locked schema block (each required by a scored *behavioral* requirement that has nowhere else to live — logged in full in `ASSUMPTIONS.md`):

- `readings.anomaly` — a reading whose `ts_device` is earlier than the zone's last-applied reading is stored (audit trail) and flagged, but never rewrites current state (TC18c).
- `zone_transitions.reason` — carries the admin's free-text justification required by "manual state set **with reason**" (`POST /api/admin/override`).
- `incidents.hazard` — TC14 requires the incident log to show and filter by *hazard type*. Set once at incident open from the fusion contributions at that instant: every hazard contributing ≥ 10 points, largest first (`"fire"`, `"fire+water"`, ...), or `"manual"` for an admin-override-opened incident, so the log also shows *what kind* of event each incident was, not just that one happened.

Race safety: `acknowledgments.incident_id` is `UNIQUE`. Ack is implemented as an insert that relies on the DB constraint, not an application-level check-then-write — see `backend/ack_service.py`. Two concurrent ack requests for the same incident: exactly one succeeds (200), the other gets 409, enforced by SQLite itself (TC7b).

7. API Reference

Base URL: `http://<host>:8000`. Auth is either a zone API key (`x-Zone-Key` header, issued per zone by `scripts/init_db.py`) or a bearer user token (`Authorization: Bearer <token>`, issued by `/api/auth/login`).

Method	Path	Auth	Purpose
POST	/api/ingest	X-Zone-Key	Raw sensor readings in; validated, deduped by (zone_id, seq)
GET	/api/zones/status	bearer	All zones' current state + priority queue in one call
GET	/api/zones/{id}/trend	bearer	Bonus 2 — slope/rising flag for one zone
GET	/api/incidents? from=&to=&zone=&status=&hazard=	bearer	History, filterable (incl. hazard type, TC14)
GET	/api/incidents/{id}	bearer	Full transition timeline for one incident (TC14)
POST	/api/incidents/{id}/ack	staff/admin	Exactly-once acknowledgment
POST	/api/admin/override	admin only	Manual state set with reason, cause='manual'
GET	/api/admin/health	admin only	System health
POST	/api/auth/login	—	Returns bearer token + role
GET	/api/commands/{zone_id}	X-Zone-Key	Zone node polls actuation commands
WS	/ws?token=	bearer (query param)	Pushes snapshot / state_change / incident_ack / periodic_snapshot
POST	/api/report	staff/admin	Bonus 4 — NL incident reporting
POST	/api/camera	X-Zone-Key	Bonus 1 — camera occupancy check-in
GET	/api/ping	—	Liveness check

POST /api/ingest

```
// Request
{
  "seq": 42,
  "fire": 0,
  "gas_norm": 0.35,
  "water_norm": 0.10,
  "occupancy": 1,
  "ts_device": "2026-07-25T09:00:00Z",
  "uptime_ms": 65000
}
// 200 Response
{ "duplicate": false, "anomaly": false, "state": "WARNING", "risk_score": 27.5 }
```

gas_norm / water_norm outside [0.0, 1.0], fire / occupancy not 0/1 → 422 with a field-level error body (TC6b, TC23f). Unregistered X-Zone-Key → 401 (TC10a). A field that is null /omitted marks that sensor offline for the reading, contributing 0 to the score — never fabricated as SAFE.

GET /api/zones/status

```
// 200 Response (abridged to one zone)
{
  "zones": [
    {
      "zone_id": 2, "name": "Server Room", "state": "CRITICAL",
      "risk_score": 65.0, "offline": false,
      "sensors": [
        {"type": "fire", "status": "online", "value": 1},
        {"type": "water", "status": "online", "value": 1.0},
        {"type": "pir", "status": "online", "value": 1}
      ],
      "last_reading_at": "2026-07-25T09:00:03Z",
      "open_incident_id": 7, "incident_status": "open",
      "trend": {"scores": [25.0, 25.0, 25.0, 65.0], "slope": 10.0, "rising": false},
      "predicted_risk": {"p_critical": 0.87, "likely": true,
        "model": "logistic_regression", "horizon_seconds": 120},
      "camera": {"occupied": true, "confidence": 0.9, "fresh": true,
        "agrees_with_pir": true}
    }
  ],
  "priority_queue": [
    {"zone_id": 2, "zone_name": "Server Room", "risk_score": 65.0,
      "occupied": true, "unacked_seconds": 26.0, "priority": 82.0,
      "justification": "Risk 65 + Occupied +15 + unacked 26s +2 = 82"}
  ],
  "server_time": "2026-07-25T09:00:04Z"
}
```

No token → 401 (TC10b).

GET /api/zones/{id}/trend

```
{ "zone_id": 1, "scores": [10, 15, 22, 30, 41], "slope": 7.8, "rising": true }
```

GET /api/incidents?zone=2&status=open&hazard=water

```
[
  {
    "id": 7, "zone_id": 2, "zone_name": "Server Room",
    "opened_at": "2026-07-25T09:00:00Z", "peak_risk": 72.5,
    "status": "open", "resolved_at": null,
    "hazard": "fire+water", "ack": null
  }
]
```


`hazard=water` matches both `"water"` and compound labels like `"fire+water"`; `hazard>manual` isolates override-opened incidents.

GET `/api/incidents/{id}`

```
{
  "incident": { "...": "IncidentOut, as above" },
  "transitions": [
    {"id": 40, "from_state": "SAFE", "to_state": "WARNING", "risk_score": 42.0,
      "cause": "sensor", "reason": null, "ts": "2026-07-25T08:59:57Z"},
    {"id": 41, "from_state": "WARNING", "to_state": "CRITICAL", "risk_score": 72.5,
      "cause": "sensor", "reason": null, "ts": "2026-07-25T09:00:00Z"}
  ]
}
```

Unknown id → 404.

POST `/api/incidents/{id}/ack`

```
// 200 Response
{ "incident_id": 7, "acked_by": 3, "ts": "2026-07-25T09:00:10Z" }
```

Unknown id → 404. Already acked (loser of a concurrent race, or a second click) → 409 (TC7b).

POST `/api/admin/override`

```
// Request
{ "zone_id": 2, "target_state": "CRITICAL", "reason": "fire drill" }
// 200 Response
{ "zone_id": 2, "state": "CRITICAL", "transitioned": true }
```

Non-admin bearer token → 403, enforced inside the handler via `Depends(require_admin)` — not just hidden in the UI (TC13). Demoed both ways: staff sees no override panel, and a direct `curl` with a staff token still gets 403.

GET `/api/admin/health`

```
{
  "status": "ok", "server_time": "2026-07-25T09:00:04Z",
  "zone_count": 3, "zones_online": 2, "reading_count": 10432,
  "open_incidents": 1
}
```

POST /api/auth/login

```
// Request
{ "username": "staff1", "password": "staff123" }
// 200 Response
{ "token": "Ab12...", "role": "staff", "username": "staff1" }
```

Bad credentials → 401.

GET /api/commands/{zone_id}

```
{
  "zone_id": 2, "state": "CRITICAL", "buzzer": true, "relay": true,
  "led": "red", "ts": "2026-07-25T09:00:00Z", "cause": "sensor"
}
```

Derived solely from the zone's latest `zone_transitions` row — WARNING gets `led: "yellow"`, `buzzer / relay` false; only CRITICAL sets both true (TC5). A zone key may only poll its own `zone_id` (403 otherwise).

WS /ws?token=<bearer>

```
// Server -> client, on connect
{ "event": "snapshot", "zones": [ "... same shape as /api/zones/status ..." ], "priority_queue": [...], "server_time": "..." }
// Server -> client, on a real transition
{ "event": "state_change", "zones": [...], "priority_queue": [...], "server_time": "..." }
```

Missing/invalid token → connection closed with code 4401. Client reconnect always triggers a full `GET /api/zones/status` refetch — no delta-application, no stale data (TC23d).

POST /api/report (Bonus 4)

```
// Request
{ "text": "small water leak near the racks in Server Room" }
// 200 Response
{
  "understood": { "zone_id": 2, "zone_name": "Server Room", "hazard_type": "water",
                  "severity": 0.3, "source": "fallback" },
  "message": "Logged as a water report for Server Room (severity 0.30, parsed via fallback). Advisory only: ..."
}
// 422 (unknown zone/hazard)
{ "detail": "unknown zone 'Mars Base' - must be one of ['Data Science Lab', 'IoT Lab', 'Server Room']" }
```

See §10 for the full safety contract.

POST /api/camera (Bonus 1)

```
// Request (X-Zone-Key header, camera node may only report its own zone)
{ "zone_id": 1, "occupied": true, "confidence": 0.9,
  "ts_device": "2026-07-25T09:00:00Z" }

// 200 Response
{ "zone_id": 1, "camera_occupied": true, "pir_occupied": false,
  "agrees_with_pir": false, "ts": "2026-07-25T09:00:00.4Z" }
```

See §11 for how the cross-check feeds the priority ranking (and nothing else).

8. RBAC

Enforced by FastAPI dependencies at the top of every handler (backend/security.py::require_admin / require_staff_or_admin), not only hidden in the UI:

Role	Can see	Can do
staff	zones, priority queue, incident timeline	acknowledge incidents, submit NL reports
admin	everything staff sees, + system health	+ manual override, + everything staff can do

Demonstrated in the video both ways: the override panel and health tab are absent from a staff login (UI), and a staff bearer token against POST /api/admin/override via curl still returns 403 (API-level, not just a hidden button) — TC13.

Accessibility consideration (TC16b) — never color alone. Every zone state is triple-encoded: color, icon, and text label (✓ SAFE , ▲ WARNING , × CRITICAL), so the dashboard remains fully readable for color-blind operators. The same principle carries through the rest of the UI: OFFLINE is a labeled badge (not a grey tint), new-CRITICAL alerts pair the visual toast with a distinct audio cue (a second sensory channel, not just more pixels), the priority queue marks its top entry with an explicit #1 rank and a text banner rather than only a pulse animation, and the Predicted Risk chip is distinguished from live state by shape and wording (dashed border, italic, "advisory" label), not color alone.

9. Bonus 2 — Short-Term Risk Trend

backend/trend.py : least-squares slope over each zone's last (up to) 8 computed risk scores. rising is set only when the slope exceeds a fixed threshold and the zone is currently WARNING — a CRITICAL zone already has the priority queue as its escalation path, and a SAFE zone drifting slightly isn't yet actionable. Embedded directly in the same canonical zones/status snapshot the rest of the dashboard reads (rule 2 — never a separately-pollled, potentially-inconsistent number). Rendered as an inline SVG sparkline plus a ↗ trending toward CRITICAL chip.

10. Bonus 4 — NL Incident Reporting

Pipeline: free text → LLM call (DeepSeek, OpenAI-compatible chat/completions , LLM_API_KEY / LLM_BASE_URL / LLM_MODEL in .env) → if no key configured, or the call errors/times out (8 s cap) → offline keyword/regex parser (backend/nl_report.py::fallback_parse) → **the same deterministic validation**

gate either way: zone must match a real zone name, hazard must be in `{fire, gas, water}`, severity is clamped (not rejected) to `[0, 1]`. A hallucinated zone name from the LLM is rejected exactly like a fallback miss would be — the gate doesn't trust the source.

Safety statement (hard rule, CLAUDE.md #4 — predicted/ML/NL-derived values never trigger actuation): a validated report only ever calls `ZoneRuntime.add_advisory_report()`, which feeds `ZoneRuntime.advisory_boost()` — a value read in exactly one place, `priority.py`'s ranking loop, and only for zones that loop has *already* restricted to CRITICAL from live sensor data. There is no code path from `nl_report.py` into `fusion.py`, `state_machine.py`, `zone_transitions`, or `/api/commands`. It can re-order which already-critical zone responders see first; it cannot make a SAFE zone appear CRITICAL, cannot change a risk score, and cannot sound a buzzer.

The advisory term decays linearly to 0 over 10 minutes and is capped at +10 total per zone (same cap shape as the unacked-time term), so a single loud report can't permanently dominate the queue, and stacking many reports on one zone doesn't runaway past +10.

11. Bonus 1 — Camera-Based Occupancy Check

Hardware honesty first: the case specifies an ESP32-CAM, which exists only as physical hardware — Wokwi does not simulate camera frames, so on Track B there is no camera to wire. We implemented the closest legitimate Track-B equivalent: `sim/camera_node.py` runs **real frame-difference motion detection** (the exact algorithm the case suggests — consecutive grayscale frames differenced, changed-pixel ratio smoothed over 5 frames, threshold at 2%) against a **laptop webcam or a video file** via OpenCV, standing in for the ESP32-CAM the same way the Wokwi potentiometer stands in for a water-level sensor. A `--synthetic` mode (scripted occupied/empty pattern, no OpenCV) exists so the *integration* can be demonstrated on any machine — the video narration states which source is on screen.

Cross-check semantics (evaluation: "correctness of the cross-check against PIR"): the camera node POSTs `{occupied, confidence}` to `/api/camera` (zone-key auth, own-zone only) once a second. The backend compares it against the PIR's debounced state and surfaces the verdict on the zone card (`CAM occupied/empty`, highlighted when it disagrees with PIR) and in every `/api/camera` response (`agrees_with_pir`).

Priority-ranking integration (evaluation: "integration into the priority ranking"): the case's stated purpose is to "reduce false 'zone is empty' readings that would otherwise skew the priority ranking" — so that is precisely, and only, where it acts. The ranking's occupancy factor is `PIR OR (fresh camera detection)`: a dead or blocked PIR can no longer demote an occupied CRITICAL zone below an empty one. When the camera is the rescuing source, the justification line says so explicitly — `occupied (camera) +15`. Camera readings go stale after 5 s (a silent camera is never trusted), and the zone's **own risk score still uses PIR alone** — the locked fusion formula is untouched, enforced by `tests/test_camera.py::test_effective_occupied_never_touches_risk_score`.

12. Bonus 3 — ML Risk Prediction

(a) Training data — synthetic, stated plainly. No real campus data exists for this system, so `ml/train.py` generates it: 400 episodes of simulated sensor behavior whose feature/label pairs are computed by the *real* backend code (`backend.fusion.risk_score`, `backend.trend.compute_slope`) — not a reimplementing that could drift. One stated modeling assumption matters: with the locked weights, CRITICAL (≥ 65) is unreachable without fire (gas 25 + water 25 + occupancy 10 caps at 60), so "predict CRITICAL" is really "predict ignition" — and a spark from nowhere is unpredictable by definition. The synthetic world therefore models the *predictable* ignition path:

smoldering (fumes rising for ~2 minutes before open flame — the case's own soldering-flux and battery-off-gassing scenarios), while keeping no-precursor sudden fires in the data so the model stays honest about what it cannot foresee.

(b) Model and why. Logistic regression (scikit-learn) over 8 features: fire level, gas/water levels, occupancy, gas/water/score slopes (least- squares over the last 8 samples), and current risk score. Chosen because the case calls it "a completely reasonable choice", it is auditable (8 readable coefficients), and it exports to plain JSON — the backend serves `m1/model.json` in pure Python (`sigmoid(w·x + b)` , `backend/prediction.py`) with no sklearn, numpy, or pickle at runtime.

(c) Validation — actually reported. Held-out test set (25%, stratified; alert threshold chosen on a separate validation split, F1-maximal):

Metric	Value
Accuracy	0.837
Precision	0.449
Recall	0.863
Decision threshold	0.498
Base positive rate	0.137

These are honest numbers, not tuned-for-show ones: the high-recall / moderate-precision operating point is deliberate for an early-warning advisory — it catches ~86% of real onsets 1–2 minutes ahead, at the cost of also flagging some rising-but-self-resolving ramps. The same metrics ship inside `m1/model.json` , so the documented numbers and the deployed artifact cannot drift apart.

(d) Dashboard presentation. A `Predicted Risk N% (ML, advisory)` chip on the zone card — dashed border, italic, violet, deliberately unlike any live-state styling — shown only when a model artifact is present, the zone has ≥ 8 readings of history, and the zone is **not** already CRITICAL (a live alarm supersedes any forecast). It is a separate field (`predicted_risk`) in the API payload, never blended into `risk_score` .

(e) Safety statement — enforced, not promised. The predicted value never triggers the relay, buzzer, a state change, or an incident. This is structural: `backend/prediction.py` is imported by the display-payload builder (`status_service.py`) only; there is no import path from it into `pipeline.py` , `state_machine.py` , `fusion.py` , or `routers/commands.py` . `tests/test_prediction.py::test_no_actuation_module_imports_prediction` scans those sources and fails the build if anyone ever wires prediction toward actuation.

13. Circuit Diagrams Per Zone

`firmware/zone_node.ino` is one sketch; only the `PER-ZONE CONFIG` block changes per zone (`ZONE_NAME` , `ZONE_ID` , `ZONE_API_KEY` , `ZONE_HAS_GAS`). Each implemented zone has its own committed wiring file (TC26):

Zone	Diagram file	Sensors wired
IoT Lab	<code>firmware/diagram.json</code>	fire + gas + water + PIR (all actuators)
Server Room	<code>firmware/diagram_server_room.json</code>	fire + water + PIR (all actuators)
Data Science Lab	<code>firmware/diagram_data_science_lab.json</code>	fire + water + PIR (all actuators)

The two no-gas variants are the IoT Lab diagram minus the `gas1` part and its 3 connections (generated programmatically, not hand-edited — the shared wiring is byte-identical). In each zone's Wokwi project, rename that zone's file to `diagram.json` and set `ZONE_HAS_GAS` accordingly.

> **Insert Wokwi screenshots here before exporting the final PDF** — one per zone, taken from the actual simulator after the human smoke-test (`firmware/README.md` "Validation status"). Not yet available from this build environment (no Wokwi runtime here); the pin table below is complete and sufficient to rebuild the circuit without the screenshot, but the screenshot is still required by the rubric.

Signal	GPIO	Wokwi part (substitution reason)
Flame (digital in)	33, INPUT_PULLDOWN	wokwi-pushbutton — no native flame/IR part; press-and-hold = sustained flame
Gas AO (analog in, IoT Lab only)	34	wokwi-gas-sensor — native MQ-2 simulation
Water SIG (analog in)	35	wokwi-potentiometer — no native water-level part; knob = simulated level
PIR OUT (digital in)	32	wokwi-pir-motion-sensor — native part
LED green	26 (220 Ω resistor)	SAFE indicator
LED yellow	27 (220 Ω resistor)	WARNING indicator
LED red	14 (220 Ω resistor)	CRITICAL indicator
Buzzer	25	CRITICAL only
Relay IN	4	CRITICAL only

GPIO 0/2/5/12/15 (ESP32 strapping pins) are deliberately unused by any output — see `firmware/README.md` for the reasoning (Wokwi won't flag a strapping-pin misuse; real hardware will refuse to boot correctly).

14. Backup & Recovery

```
scripts/backup.sh # sqlite3 .backup - online, WAL-safe, no downtime
```

Writes `backups/scsrg_<timestamp>.db`. **Restore:** stop the server, then `cp backups/scsrg_<timestamp>.db data/scsrg.db`, then restart (which rebuilds all in-memory zone state from that DB per rule 6). **Gap:** anything written between the last backup and a crash is lost. This is an accepted, documented gap — high-volume/low-value `readings` rows are the likely casualty, while `incidents` / `acknowledgments` (the rows that actually matter for audit/compliance) are written synchronously well before a typical backup cadence (e.g. hourly cron) would run.

15. Retention & Access Policy

Raw `readings` older than 90 days are intended to be summarized (e.g. hourly min/max/avg per zone) and dropped — the fusion formula only ever needs the last few seconds of readings per zone, so 90-day raw retention is already generous for anything except forensic replay, and `readings` is by far the fastest-growing table (one row per zone per 750 ms). `incidents`, `zone_transitions`, and `acknowledgments` are kept indefinitely — they're the compliance/audit record, orders of magnitude smaller in volume, and exactly the rows TC14's timeline and TC19's

history queries depend on. **Only admins** query raw `readings` history directly (ties to the RBAC model in §8 — `staff` can see current zone state and incident timelines, but bulk raw-sensor export is an admin-only capability, matching who's authorized to run `/api/admin/health` and overrides).

16. Scaling Note (TC11)

Today: one uvicorn process, one SQLite file (WAL mode), in-process `asyncio.Lock` per zone for write serialization, in-process pub/sub for WS fan-out. Verified responsive at `sim/driver.py --phantom 30` (30 concurrent phantom zones posting every 750 ms, 0 errors, ~460–810 ms p95 latency) and at 10k+ seeded readings (`sim/seed.py` , indexed incident query ~14 ms).

For a real 30+ zone campus, the single-process assumptions are exactly what would need to change, each behind an interface that's already isolated in this codebase:

- **uvicorn workers** (`--workers N`) once ingestion volume exceeds one process — requires moving the per-zone `asyncio.Lock` (currently in-process, `zone_manager.py`) to something cross-process, e.g. a Postgres row lock or advisory lock on `zones.id` .
- **SQLite → Postgres**: same SQLAlchemy models, swap `DATABASE_URL` and the WAL pragmas; `ON DELETE RESTRICT` and unique constraints translate directly.
- **MQTT ingestion** in front of `/api/ingest` for real 30+ node fleets, so nodes don't all hold open HTTP connections to one process — a lightweight bridge subscribing to MQTT and calling the same `pipeline.process_reading()` would need no change to the fusion/state logic at all.
- **WS fan-out via Redis pub/sub** once there's more than one backend process — `broadcaster.py` 's `events.bus` and `ws_manager.py` are already isolated behind a small interface (`publish / subscribe` , `broadcast`), so swapping the in-process bus for a Redis-backed one is a contained change, not a rewrite of the pipeline or the dashboard contract.

17. Track Declaration

Track B (Wokwi ESP32) is primary. All three zones' circuits are defined in `firmware/diagram.json` / `firmware/zone_node.ino` and are intended to be demonstrated live in the Wokwi simulator. `sim/` (Python) is used exclusively for load generation (`--phantom 30`), protocol-level edge cases that need precise timing/malformed input (duplicate seq, malformed payload, concurrent races), and the 10k-row perf seed — never as a substitute for the Wokwi hardware demo. See `firmware/README.md` for the current validation status and the Wokwi project link(s) (added to `README.md` once smoke-tested).

18. Testing Summary

98 pytest cases (`pytest -q`) covering fusion math, flame debounce/decay, PIR hold, CRITICAL hysteresis, duplicate-seq dedup, out-of-order/anomaly flagging, incident open/resolve/re-trigger + hazard labeling, concurrent ack race, restart recovery, WS snapshot/broadcast, 401/403/404/409/422 paths, Bonus 2 trend math, Bonus 4 parsing/validation/advisory-boost decay+cap, Bonus 3 prediction math + artifact honesty + the no-actuation-import source guard, and Bonus 1 camera cross-check (freshness, PIR agreement, priority rescue, formula isolation). Every pytest-covered behavior has also been re-verified live: a real uvicorn process with a real WebSocket client (not `TestClient`), and the dashboard driven end-to-end in an actual browser. `sim/driver.py` (all scenario groups tc1–tc24) is green end-to-end against a live server. See `README.md` for exact commands.

